

Class 9 Candy Mini-Project

Katherine Quach (A18541014)

Table of contents

Background	1
Data Input	1
Exploratory Analysis	4
Exploring the Correlation Structure	16
Principal Component Analysis (PCA)	17

Background

In this mini-project, you will explore FiveThirtyEight's Halloween Candy dataset.

We will use lots of **ggplot** some basic stats, correlation analysis and PCA to make sense of the landscape of US candy - something hopefully more relateable than the proteomics and transcriptomics work that we will use these methods on throughout the rest of the course.

Data Input

Our data is a CSV file so we use `read.csv()`

```
candy <- read.csv("candy-data.csv", row.names = 1)
head(candy)
```

	chocolate	fruity	caramel	peanut	yalmond	nougat	crisped	rice	wafer
100 Grand	1	0	1		0	0			1
3 Musketeers	1	0	0		0	1			0
One dime	0	0	0		0	0			0
One quarter	0	0	0		0	0			0
Air Heads	0	1	0		0	0			0
Almond Joy	1	0	0		1	0			0

	hard	bar	pluribus	sugarpercent	pricepercent	winpercent
100 Grand	0	1	0	0.732	0.860	66.97173
3 Musketeers	0	1	0	0.604	0.511	67.60294
One dime	0	0	0	0.011	0.116	32.26109
One quarter	0	0	0	0.011	0.511	46.11650
Air Heads	0	0	0	0.906	0.511	52.34146
Almond Joy	0	1	0	0.465	0.767	50.34755

Q1. How many different candy types are in this dataset?

```
candy$fruity
```

```
[1] 0 0 0 0 1 0 0 0 0 1 0 1 1 1 1 1 1 1 0 1 1 0 0 0 1 0 0 1 1 1 0 0 0  

[39] 0 0 0 1 0 0 1 1 0 0 0 1 1 0 0 0 0 1 0 0 1 0 1 1 0 1 0 0 1 1 1 1 0 0 1 1 1 0  

[77] 0 0 1 0 1 1 1 0 0
```

```
nrow(candy)
```

```
[1] 85
```

Q2. How many fruity candy types are in the dataset?

```
sum(candy$fruity)
```

```
[1] 38
```

Q3. What is your favorite candy (other than Twix) in the dataset and what is it's winpercent value?

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from 'package:stats':
```

```
filter, lag
```

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
candy |> filter(row.names(candy) == "Sour Patch Kids") |>  
  select(winpercent)
```

```
      winpercent  
Sour Patch Kids      59.864
```

Q4. What is the winpercent value for “Kit Kat”?

```
candy["Kit Kat", ]$winpercent
```

```
[1] 76.7686
```

I prefer this format because it's shorter than loading the library()

Q5. What is the winpercent value for “Tootsie Roll Snack Bars”?

```
candy["Tootsie Roll Snack Bars", ]$winpercent
```

```
[1] 49.6535
```

Q6. Is there any variable/column that looks to be on a different scale to the majority of the other columns in the dataset?

Yes. The winpercent column has values that are beyond the 0-1 scale that the other columns have.

```
library("skimr")  
skim(candy)
```

Table 1: Data summary

Name	candy
Number of rows	85
Number of columns	12
Column type frequency:	
numeric	12
Group variables	None

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
chocolate	0	1	0.44	0.50	0.00	0.00	0.00	1.00	1.00	
fruity	0	1	0.45	0.50	0.00	0.00	0.00	1.00	1.00	
caramel	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
peanutyalymond	0	1	0.16	0.37	0.00	0.00	0.00	0.00	1.00	
nougat	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
crispedrice-wafer	0	1	0.08	0.28	0.00	0.00	0.00	0.00	1.00	
hard	0	1	0.18	0.38	0.00	0.00	0.00	0.00	1.00	
bar	0	1	0.25	0.43	0.00	0.00	0.00	0.00	1.00	
pluribus	0	1	0.52	0.50	0.00	0.00	1.00	1.00	1.00	
sugarpercent	0	1	0.48	0.28	0.01	0.22	0.47	0.73	0.99	
pricepercent	0	1	0.47	0.29	0.01	0.26	0.47	0.65	0.98	
winpercent	0	1	50.32	14.71	22.45	39.14	47.83	59.86	84.18	

Q7. What do you think a zero and one represent for the `candy$chocolate` column?

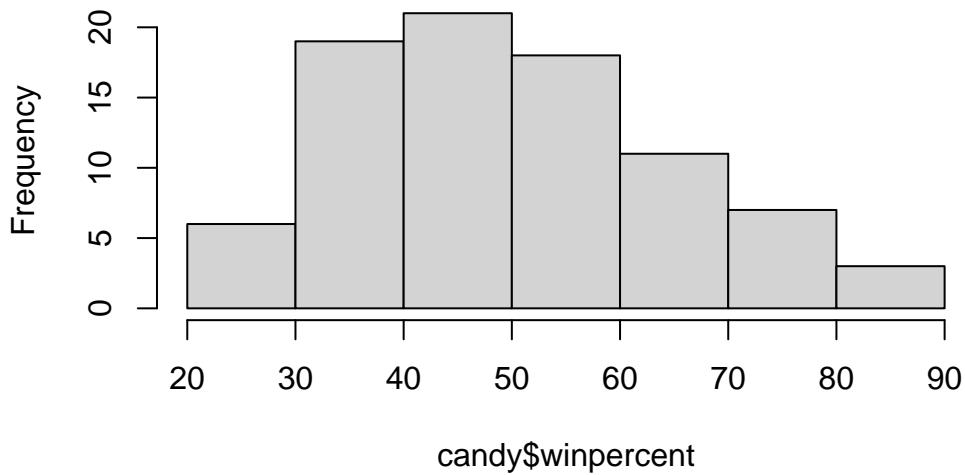
In the `candy$chocolate` column, the 0 represents FALSE, meaning that the candy does NOT contain chocolate. The 1 in the column represents TRUE, hence the candy does contain chocolate.

Exploratory Analysis

Q8. Plot a histogram of winpercent values

```
hist(candy$winpercent)
```

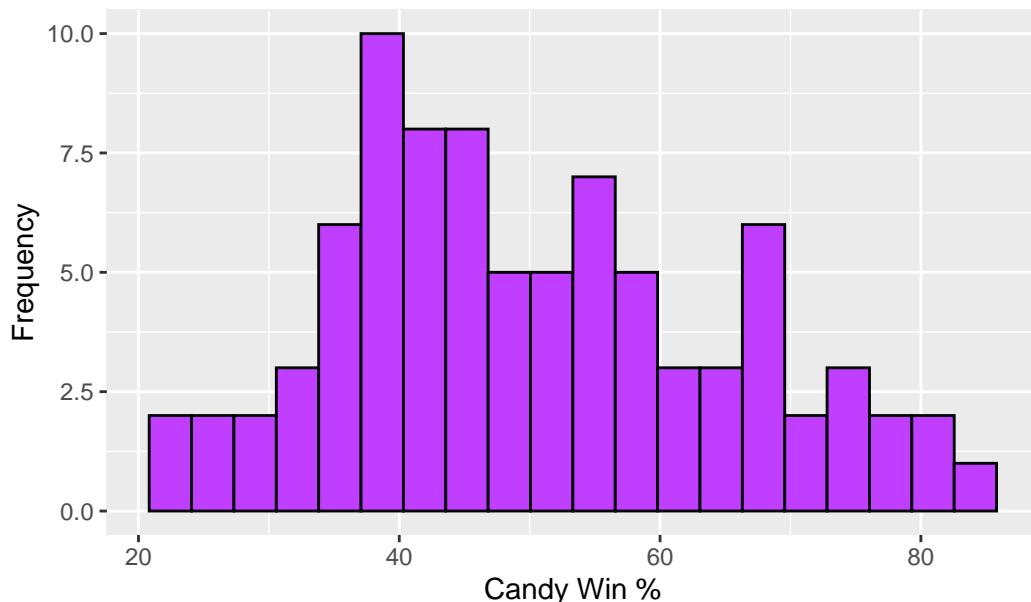
Histogram of candy\$winpercent



```
library(ggplot2)

ggplot(candy, aes(x = winpercent)) +
  geom_histogram(bins = 20, fill = "#BF3EFF", color = "black") +
  labs(
    x = "Candy Win %",
    y = "Frequency",
    title = "Distribution of Candy Win Percentage")
```

Distribution of Candy Win Percentage



Q9. Is the distribution of winpercent values symmetrical?

The distribution of winpercent values is not symmetrical

Q10. Is the center of the distribution above or below 50%?

The center of the distribution is above 50%

```
mean(candy$winpercent)
```

[1] 50.31676

```
summary(candy$winpercent)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
22.45	39.14	47.83	50.32	59.86	84.18

Q11. On average is chocolate candy higher or lower ranked than fruit candy?

1. Find all chocolate candy
2. Get their winpercent values
3. Find the mean
4. Find all fruit candy

5. Get their winpercent values
6. Find the mean
7. Compare the two means

On average, chocolate candy is ranked higher than fruit candy

```
choc.candy <- candy[candy$chocolate == 1, ]
choc.win <- choc.candy$winpercent
mean(choc.win)
```

[1] 60.92153

```
fruit.candy <- candy[candy$fruity == 1, ]
fruit.win <- fruit.candy$winpercent
mean(fruit.win)
```

[1] 44.11974

Q12. Is this difference statistically significant?

This difference is statistically significant

```
t.test(choc.win, fruit.win)
```

Welch Two Sample t-test

```
data: choc.win and fruit.win
t = 6.2582, df = 68.882, p-value = 2.871e-08
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
11.44563 22.15795
sample estimates:
mean of x mean of y
60.92153 44.11974
```

Q13. What are the five least liked candy types in this set?

```
tail(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Snickers	1	0	1		1	1
Kit Kat	1	0	0		0	0
Twix	1	0	1		0	0
Reese's Miniatures	1	0	0		1	0
Reese's Peanut Butter cup	1	0	0		1	0
	crisped	rice	wafer	hard	bar	pluribus
Snickers	0	0	1		0	0.546
Kit Kat	1	0	1		0	0.313
Twix	1	0	1		0	0.546
Reese's Miniatures	0	0	0		0	0.034
Reese's Peanut Butter cup	0	0	0		0	0.720
	price	percent	win	percent		
Snickers	0.651	76.67378				
Kit Kat	0.511	76.76860				
Twix	0.906	81.64291				
Reese's Miniatures	0.279	81.86626				
Reese's Peanut Butter cup	0.651	84.18029				

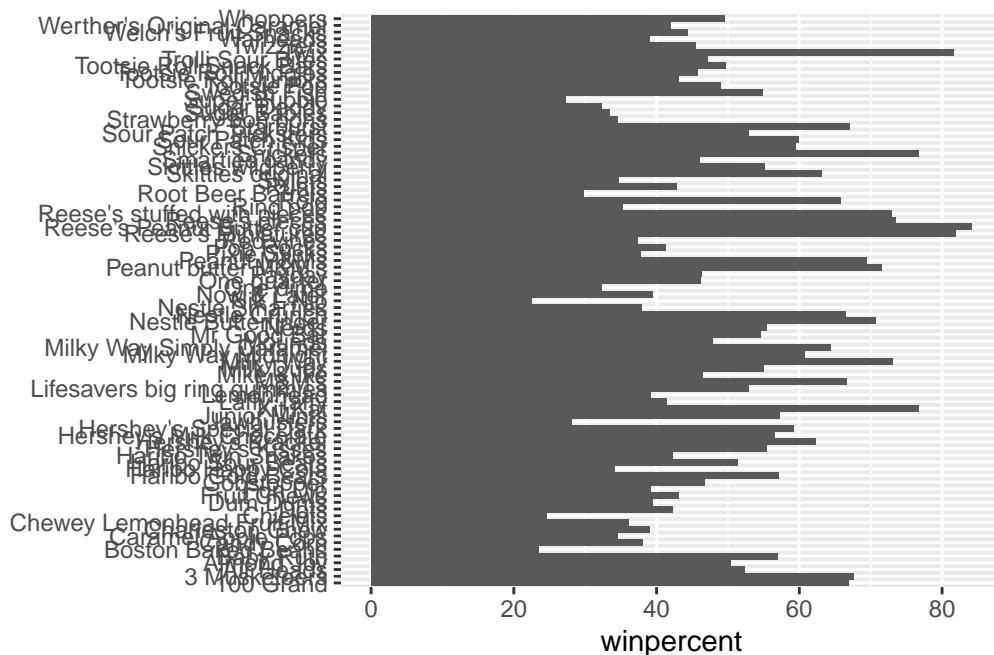
Q14. What are the top 5 all time favorite candy types out of this set?

```
head(candy[order(candy$winpercent),], n=5)
```

	chocolate	fruity	caramel	peanuty	almondy	nougat
Nik L Nip	0	1	0		0	0
Boston Baked Beans	0	0	0		1	0
Chiclets	0	1	0		0	0
Super Bubble	0	1	0		0	0
Jawbusters	0	1	0		0	0
	crisped	rice	wafer	hard	bar	pluribus
Nik L Nip	0	0	0		1	0.197
Boston Baked Beans	0	0	0		1	0.313
Chiclets	0	0	0		1	0.046
Super Bubble	0	0	0		0	0.162
Jawbusters	0	1	0		1	0.093
	win	percent				
Nik L Nip	22.44534					
Boston Baked Beans	23.41782					
Chiclets	24.52499					
Super Bubble	27.30386					
Jawbusters	28.12744					

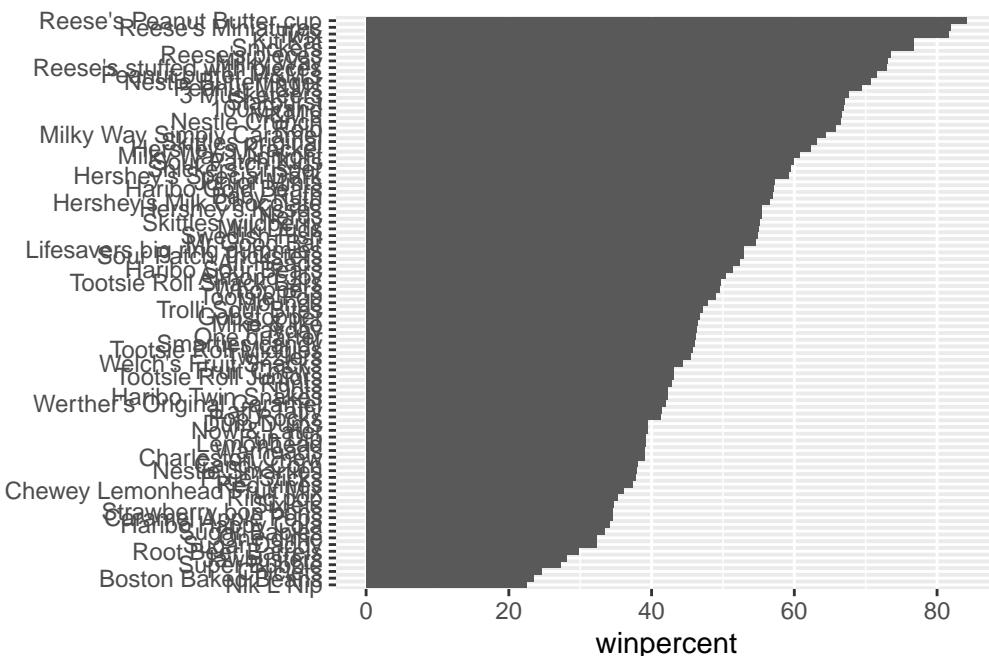
Q15. Make a first barplot of candy ranking based on winpercent values.

```
ggplot(candy) +
  aes(winpercent, rownames(candy)) +
  geom_col() +
  ylab("")
```

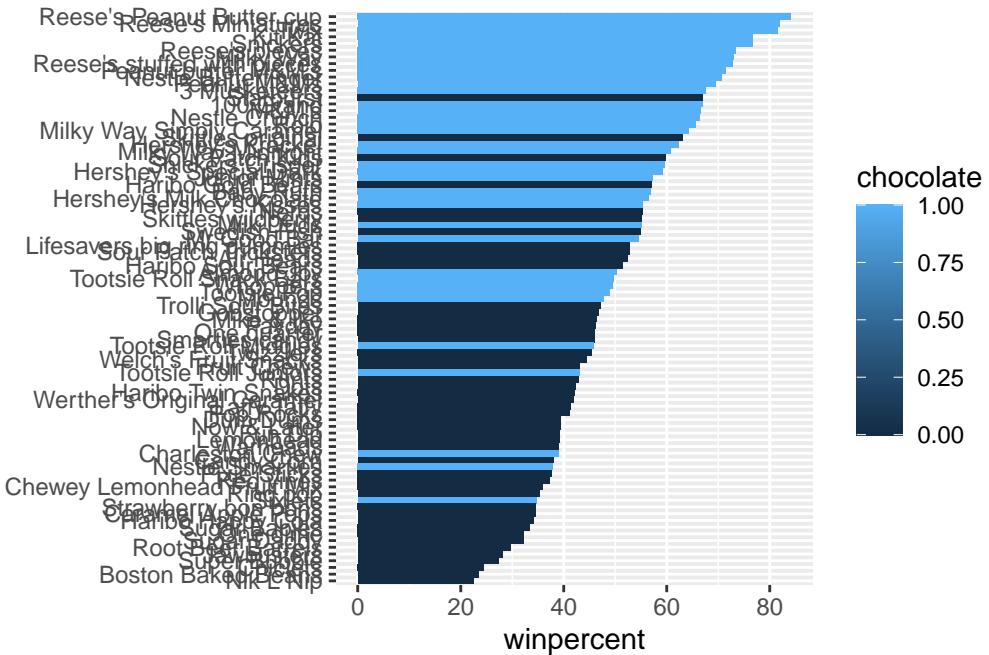


Q16. This is quite ugly, use the reorder() function to get the bars sorted by winpercent?

```
ggplot(candy) +
  aes(winpercent,
      reorder(rownames(candy), winpercent)) +
  geom_col() +
  ylab("")
```



```
ggplot(candy) +
  aes(winpercent,
      reorder(rownames(candy), winpercent),
      fill = chocolate) +
  geom_col() +
  ylab("")
```

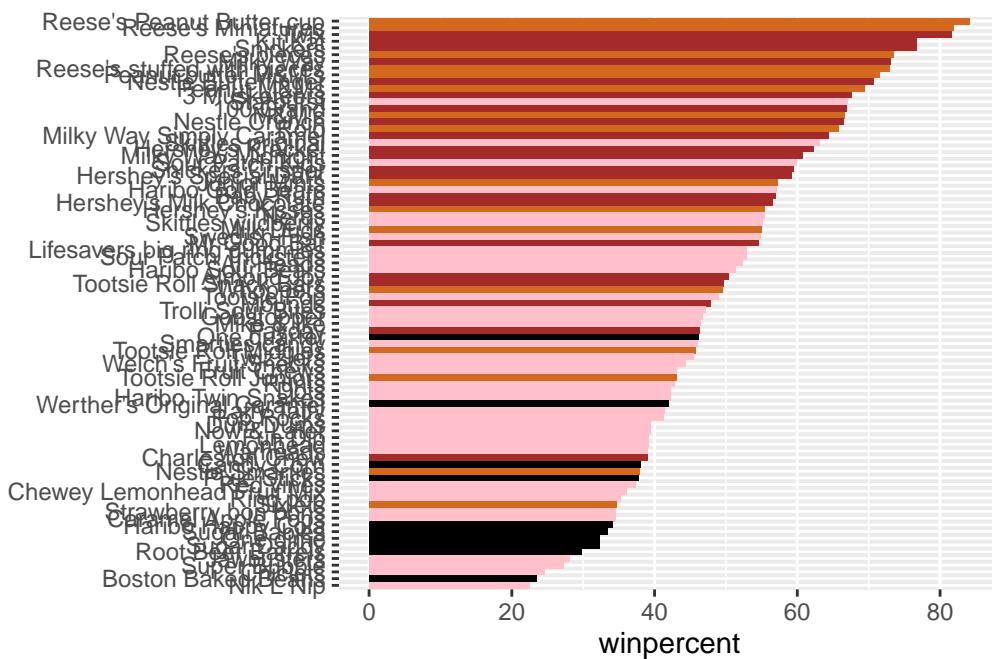


We need a custom color vector

```
my_cols <- rep("black", nrow(candy))
my_cols[candy$chocolate == 1] <- "chocolate"
my_cols[candy$bar == 1] <- "brown"
my_cols[candy$fruity == 1] <- "pink"
my_cols
```

```
[1] "brown"      "brown"       "black"        "black"        "pink"        "brown" 
[7] "brown"      "black"       "black"        "pink"        "brown"       "pink"  
[13] "pink"       "pink"        "pink"         "pink"        "pink"        "pink"  
[19] "pink"       "black"       "pink"         "pink"        "pink"        "brown" 
[25] "brown"      "brown"       "pink"         "chocolate"   "brown"       "pink"  
[31] "pink"       "pink"        "chocolate"    "chocolate"   "pink"        "chocolate" 
[37] "brown"      "brown"       "brown"        "brown"       "brown"       "pink"  
[43] "brown"      "brown"       "pink"         "pink"        "brown"       "chocolate" 
[49] "black"       "pink"        "pink"         "chocolate"   "chocolate"   "chocolate" 
[55] "chocolate"   "pink"        "chocolate"    "black"       "pink"        "chocolate" 
[61] "pink"       "pink"        "chocolate"    "pink"        "brown"       "brown"  
[67] "pink"       "pink"        "pink"         "pink"        "black"       "black"  
[73] "pink"       "pink"        "pink"         "chocolate"   "chocolate"   "brown"  
[79] "pink"       "brown"       "pink"         "pink"        "pink"        "black" 
[85] "chocolate"
```

```
ggplot(candy) +
  aes(winpercent,
      reorder(rownames(candy), winpercent)) +
  geom_col(fill = my_cols) +
  ylab("")
```



Q17. What is the worst ranked chocolate candy?

The worst ranked chocolate candy is Sixlets

```
rownames(candy)[candy$chocolate == 1][which.min(candy$winpercent[candy$chocolate == 1])]
```

```
[1] "Sixlets"
```

Q18. What is the best ranked fruity candy?

The best ranked fruity candy is Starburst

```
rownames(candy)[candy$fruity == 1][which.max(candy$winpercent[candy$fruity == 1])]
```

```
[1] "Starburst"
```

```

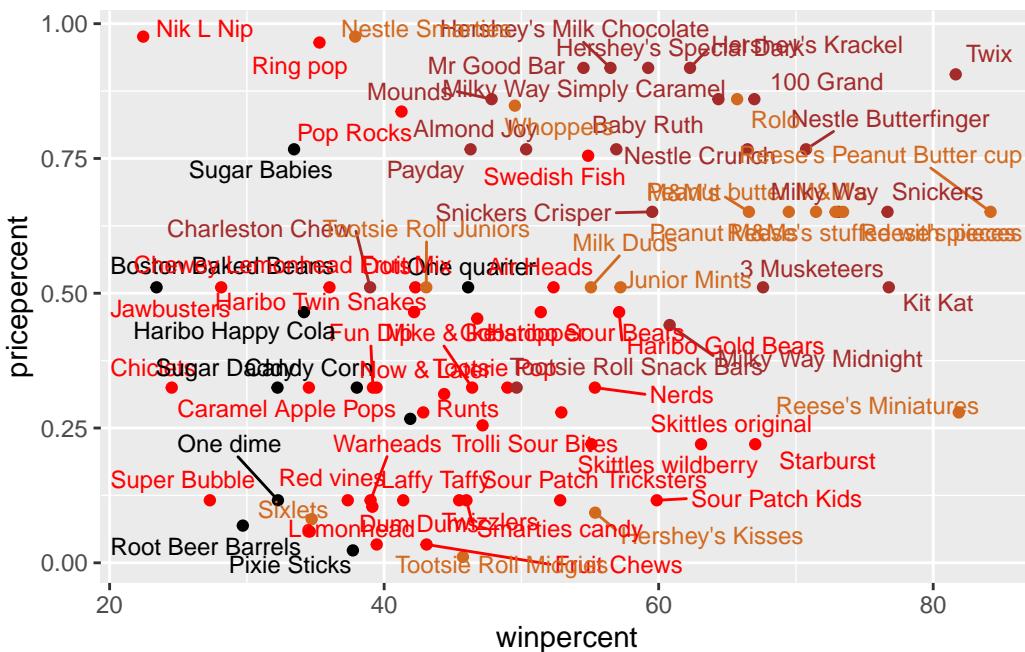
library(ggrepel)

my_cols[candy$fruity == 1] <- "red"

# How about a plot of win vs price
ggplot(candy) +
  aes(winpercent, pricepercent, label=rownames(candy)) +
  geom_point(col=my_cols) +
  geom_text_repel(col=my_cols, size=3.3, max.overlaps = 17)

```

Warning: ggrepel: 4 unlabeled data points (too many overlaps). Consider increasing max.overlaps



Q19. Which candy type is the highest ranked in terms of winpercent for the least money - i.e. offers the most bang for your buck?

The candy that's highest ranked in terms of winpercent for the least money is chocolate candy

```

min_price <- min(candy$pricepercent)

cheap_candies <- candy[candy$pricepercent == min_price, ]

```

```
best_row <- which.max(cheap_candies$winpercent)

colnames(cheap_candies)[best_row]
```

```
[1] "chocolate"
```

Q20. What are the top 5 most expensive candy types in the dataset and of these which is the least popular?

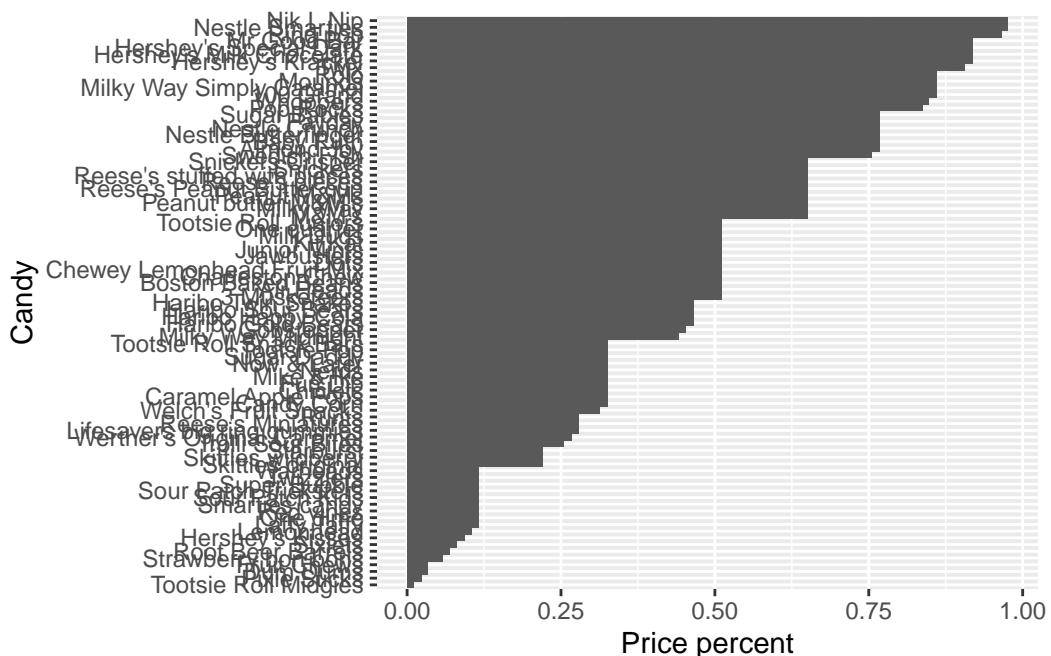
The top 5 most expensive candy types in the dataset are as follows:

```
ord <- order(candy$pricepercent, decreasing = TRUE)
head(candy[ord,c(11,12)], n=5)
```

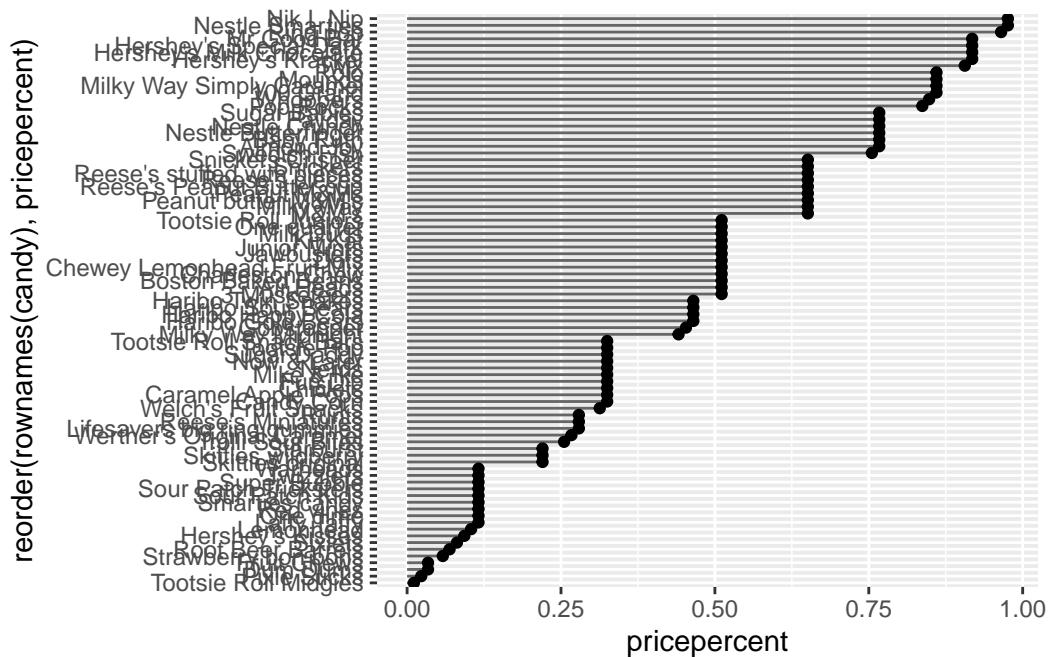
	pricepercent	winpercent
Nik L Nip	0.976	22.44534
Nestle Smarties	0.976	37.88719
Ring pop	0.965	35.29076
Hershey's Krackel	0.918	62.28448
Hershey's Milk Chocolate	0.918	56.49050

Q21. Make a barplot again with geom_col() this time using pricepercent and then improve this step by step, first ordering the x-axis by value and finally making a so called “dot chat” or “lollipop” chart by swapping geom_col() for geom_point() + geom_segment().

```
ggplot(candy,
       aes(x = reorder(rownames(candy), pricepercent),
           y = pricepercent)) +
  geom_col() +
  coord_flip() +
  labs(x = "Candy", y = "Price percent")
```



```
ggplot(candy) +
  aes(pricepercent, reorder(rownames(candy), pricepercent)) +
  geom_segment(aes(yend = reorder(rownames(candy), pricepercent),
                   xend = 0), col="gray40") +
  geom_point()
```



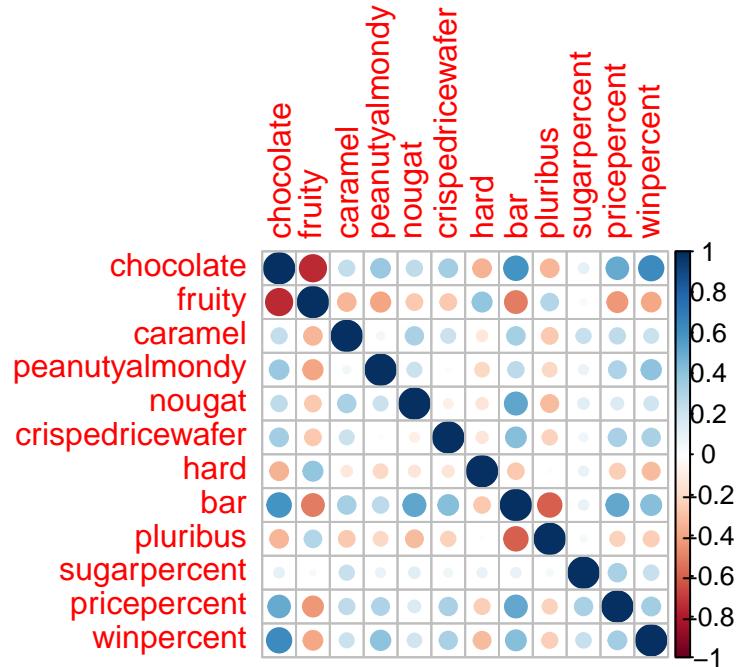
Exploring the Correlation Structure

```
cij <- cor(candy)
```

```
library(corrplot)
```

```
corrplot 0.95 loaded
```

```
corrplot(cij)
```



Q22. Examining this plot what two variables are anti-correlated (i.e. have minus values)?

The 2 variables that are anti-correlated are chocolate and fruity.

Q23. Similarly, what two variables are most positively correlated?

The 2 variables that are most positively correlated are chocolate and bar.

Principal Component Analysis (PCA)

```
candy2 <- read.csv("candy-data.csv", row.names = 1)  
  
pca <- prcomp(candy2, scale = TRUE)  
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	2.0788	1.1378	1.1092	1.07533	0.9518	0.81923	0.81530

	Proportion of Variance	0.3601	0.1079	0.1025	0.09636	0.0755	0.05593	0.05539
Cumulative Proportion	0.3601	0.4680	0.5705	0.66688	0.7424	0.79830	0.85369	
Standard deviation		0.74530	0.67824	0.62349	0.43974	0.39760		
Proportion of Variance	0.04629	0.03833	0.03239	0.01611	0.01317			
Cumulative Proportion	0.89998	0.93832	0.97071	0.98683	1.00000			

Score Plot...

```
ggplot(pca$x) +
  aes(PC1, PC2, label = row.names(pca$x)) +
  geom_point(col = my_cols) +
  geom_text_repel(max.overlaps = 30, size = 3.3, col = my_cols) +
  theme(legend.position = "none") +
  labs(title="Halloween Candy PCA Space",
       subtitle="Colored by type: chocolate bar (dark brown), chocolate other (light brown),
```

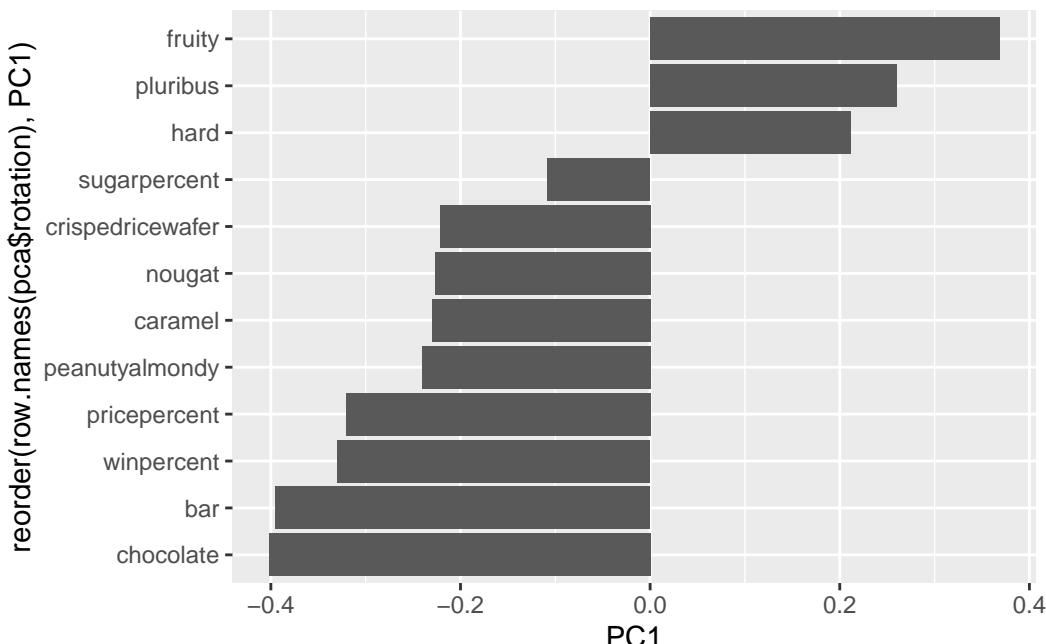
Warning: ggrepel: 4 unlabeled data points (too many overlaps). Consider increasing max.overlaps

Halloween Candy PCA Space

Colored by type: chocolate bar (dark brown), chocolate other (light brown),



```
ggplot(pca$rotation) +
  aes(PC1,
      y = reorder(row.names(pca$rotation), PC1), x = PC1) +
  geom_col()
```



Q24. Complete the code to generate the loadings plot above. What original variables are picked up strongly by PC1 in the positive direction? Do these make sense to you? Where did you see this relationship highlighted previously?

The variables that are picked up strongly by PC1 in the positive direction are “fruity”, “pluribus”, and “hard”. Yes, this does make sense. PC1 separates the fruity, hard, sugary, and non-chocolate candies from the chocolate candies. The positive loadings are hard, fruit candies, whereas the negative loadings describe chocolate bars (e.g. chocolate, bar, caramel, nougat, etc.). This relationship was previously highlighted in the correlation structure where chocolate and fruity were the most anti-correlated pair. Hard candies also tended to be non-chocolate, which was denoted by PC1. This was also evident in the PCA scatter plot where fruity candies were clustered on one side of PC1 and chocolate was clustered on the opposite side. PC1 visually separates the candy types in the same way the loading plots above explain numerically.

Q25. Based on your exploratory analysis, correlation findings, and PCA results, what combination of characteristics appears to make a “winning” candy? How do these different analyses (visualization, correlation, PCA) support or complement each other in reaching this conclusion?

To make a “winning” candy, it must be chocolate-based, include caramel, nougat, or nuts, and come bar-shaped. They are not fruity, hard, or pluribus, and skew slightly higher in pricepercent. Exploratory plots showed that these candies (e.g. chocolate bars) were more likely to be clustered at the top of winpercent, whereas fruity/hard candies were clustered farther away. The correlation matrix revealed strong positive correlations between chocolate, bar, caramel, and peanutyalmondy traits, and a strong negative correlation with fruity and hard candies. PCA loadings grouped chocolate, bar, caramel, nougat, and nuts on the opposite side of fruity/hard candies. This indicates that through visualization, correlation, and PCA analyses, chocolate bar with caramel/nougat/nuts with a slightly premium price makes a “winning” candy.